

Lab 2. Data Visualization and Graphics



Data visualizations are very important in data science. They are used as a part of **Exploratory Data Analysis (EDA)**, to familiarize yourself with data, to examine the distributions of variables, to identify outliers, and to help guide data cleaning and analysis. They are also used to communicate results to a variety of audiences, from other data scientists to customers.

In this lab, you'll learn about base plots, ggplot2, and will be briefly introduced to more advanced plotting with the applications Shiny and Plotly.

By the end of this lab, you will be able to:

- Use Base R for plotting, and identify when to do so
- Create a variety of different data visualizations using the ggplot2 package
- Explain different tools for interactive plotting in R

Lab Environment

All packages have been installed. There is no requirement for any setup.

All datasets and examples are present in `~/Desktop/R-Programming/lesson2/` folder.

The `plot()` Function

`main`

an overall title for the plot: see [title](#).

`sub`

a sub title for the plot: see [title](#).

`xlab`

a title for the x axis: see [title](#).

`ylab`

a title for the y axis: see [title](#).

`asp`


the y/x aspect ratio, see [plot.window](#).

The `plot()` function is the backbone of base plots in R. It provides capability for generic $[X-Y]$ plotting. It requires only one argument, `[x]`, which should be something to plot---a vector of numbers, one variable of a dataset, or a model object such as linear or logistic regression. You can, of course, add a second variable, `[y]`, plus an assortment of options to customize the plot, but `[x]` is the only input required for the function to run successfully.

Note:

In RStudio, sometimes the plot may be skewed or squished, as it is constrained by the size of your plot window (usually the bottom-right window, under the **Plots** tab.) You can, at any time, click the **Zoom** button and your


plot will pop out, usually larger, and give you a better look:



If we first load the `datasets` library, we gain access to a number of built-in datasets in R that will be useful for both base plotting and using `ggplot2`. To begin with, we'll use the `mtcars` dataset. `mtcars` is a very famous example dataset, and its description (accessed using `?mtcars`) is as follows:


[The data was extracted from the 1974 Motor Trend US magazine, and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973--74 models).]

Minimally, we can plot just one variable of `mtcars`, for example `mpg` or the miles per gallon of the cars. This generates a very basic plot of `mpg` on the y-axis, with `index` on the x-axis, literally corresponding to the row index of each observation, as follows:



This plot isn't very informative, but it is powerful in terms of seeing how well R can plot even when it is not installed on a particular machine. Let's add in a second variable and plot `mpg` versus `wt`:

```
plot(mtcars$wt, mtcars$mpg)
```



If we plot `mpg` versus `wt` (`[y]` vs. `[x]`), we can see a clear negative linear trend, that is, when the weight increases, the miles per gallon decreases. This isn't terribly unexpected---heavier cars will require more gas to operate and will therefore get less miles to the gallon.

You should also notice that the default axis labels are the variables exactly as input into `plot()` (so, `mtcars$wt` and `mtcars$mpg` include the dataset name and dollar sign to access each variable). There is no title by default, and the default shape is an open circle.

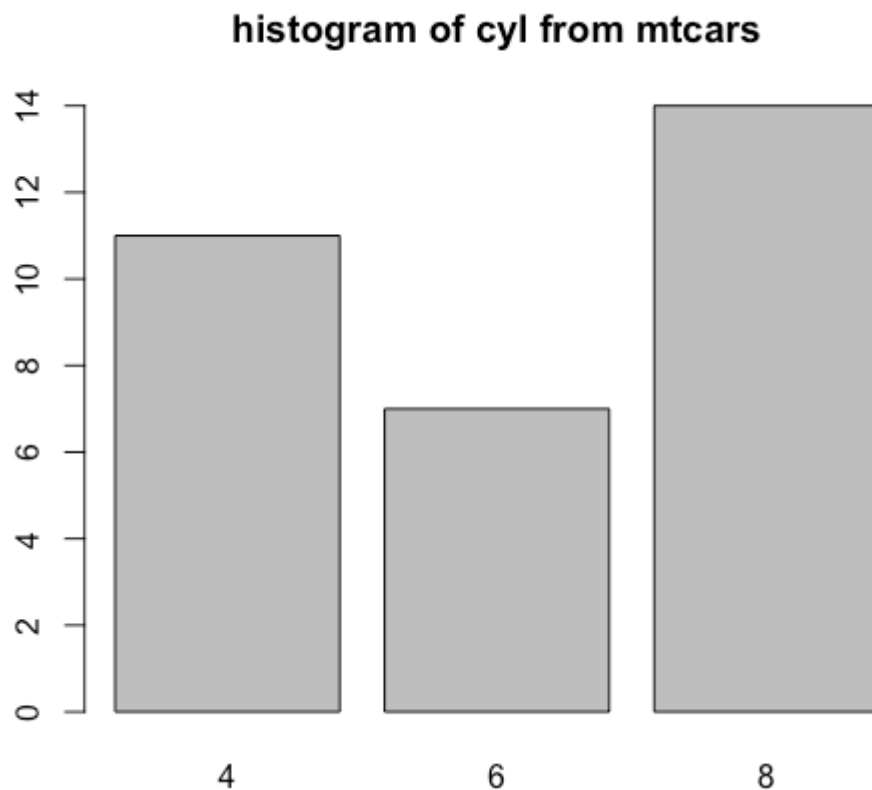
These last two plots were an example of what happens if you input variables from a dataset into `plot()`. The `plot()` function is very versatile, however, and you can input a number of different things and still create base plots. Let's discuss a few of the options in the next few subtopics.

Factor Variables

We input a few variables from `mtcars` into `plot()`, but they were continuous. What happens if, instead, we input a factor variable?

For example, the `cyl` variable in `mtcars` gives the number of cylinders each car has. If we input it as a factor variable into `plot`, we get a bar chart (histogram) by default, where each bar gives a count of how many cars have each number of cylinders:

```
plot(as.factor(mtcars$cyl))
```




Let's now create plots using factor variables and learn to differentiate between plots created with factor variables and those created without. Follow the steps given below:

1. Load the `mtcars` dataset using the `data("mtcars")` method.
2. Plot the gear variable of `mtcars` without changing it to a factor variable using `plot(mtcars$gear)` .
What kind of plot do you get?
3. Now, plot the gear variable of `mtcars` as a factor variable, as follows:


```
plot(as.factor(mtcars$gear))
```

What kind of plot is generated?

Output: The following scatterplot is the output we get when the `gear` variable is plotted without changing it to the factor variable:



The following histogram/bar chart is the output we get when the `gear` variable is plotted after changing it to the factor variable:



Model Objects

As we observed in the factor variable example, the function defaults to certain types of plots depending on the kind of data you put into it. If you were to input a linear model object, `plot()` automatically returns four helpful model diagnostic plots, including the Residuals versus Fitted and Normal Q-Q plots, which help you determine whether your model fits well. The following code demonstrates this:

```
mtcars_lm <- lm(mpg ~ wt, data = mtcars)
plot(mtcars_lm)
```

The process of generating all four of these plots is somewhat tedious, however, so instead, let's look at plotting more than one plot at a time, combined with model object plotting.

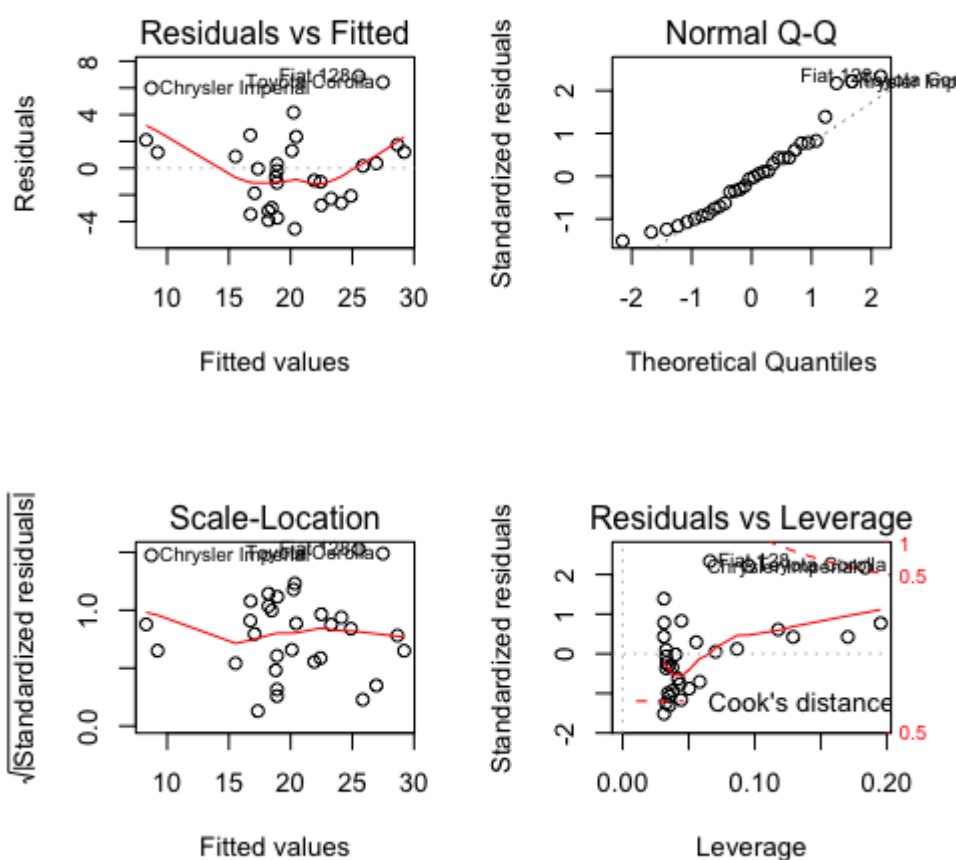
Plotting More Than One Plot at a Time

One neat feature in R is that you can plot more than one plot at a time on the same viewing window. Inside of `par()`, if we pass `mfrow = c(rows, cols)`, where `rows` is the number of rows of plots you'd like and `cols` is the number of columns of plots you'd like, you can plot a number of plots on the same screen. If we return to the `mtcars_lm()` example we just covered, we can plot all four model diagnostic plots in the same window by first running the following line of code:

```
par(mfrow = c(2,2))
```

Next, you need to execute the following code:

```
plot(mtcars_lm)
```



This resets your Global Options in RStudio. So now, every time you try and plot, it will make plots in a $[2 \times 2]$ grid. You'll need to reset back to $[*1 \times *1]$ when you're ready, using either `dev.off()` or simply `par(mfrow = c(1,1))`.

Creating and Plotting a Linear Model Object

Let's use the function to create a linear model object, and then use `par()`, `mfrow()`, and `plot()` to examine the model diagnostics. Follow the steps given below:

1. Build your own version of `mtcars_lm`, which looks at how the displacement and weight variables affect `mpg` using the following code:

```
mtcars_lm <- lm(mpg ~ disp + wt, data = mtcars)
```

2. Run the following code to enable plotting a $[*2 \times *2]$ grid of plots so that looking at model diagnostic plots is easier with the following method:

```
par(mfrow = c(2, 2))
```

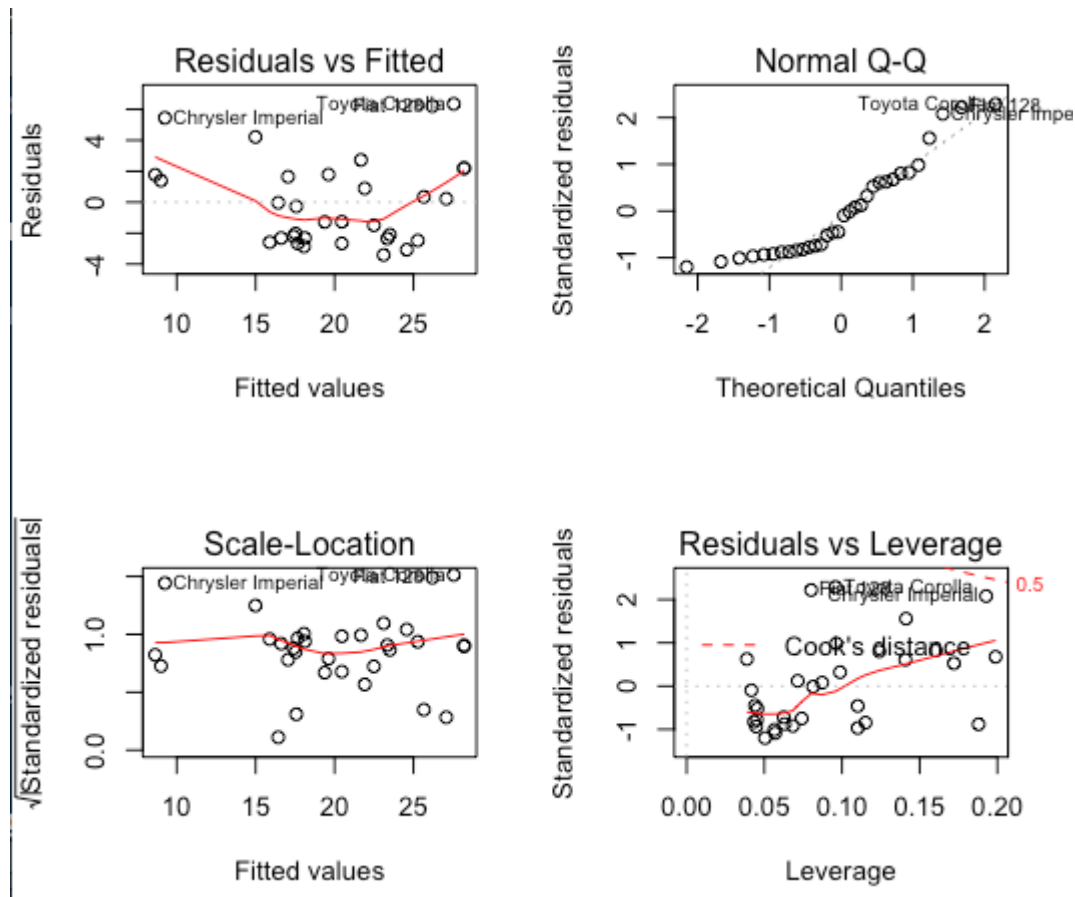
3. Plot the `mtcars_lm` variable to see the model diagnostic plots using `plot(mtcars_lm)`.

Note:

Be sure to turn the `[*2]] × [*2]` grid off. This will make the plot disappear, so be sure you're done looking at it before you run this line.

4. Turn the `[*2 ×]][*2]` grid off using `dev.off()`.

Output: The following is the output we get when we execute the `plot()` function as mentioned in [Step 3].




Titles and Axis Labels

Let's return to our `mtcars` scatterplot, add a title and subtitle, and also change the axis labels with the following code:

```
plot(mtcars$wt, mtcars$mpg,
     main = "MPG vs. Weight",
     sub = "mtcars dataset",
     xlab = "Weight",
     ylab = "MPG")
```

This adds our title and overrides the default behavior of printing axis labels, which are exactly what was input for `[x]` and `[y]`. The plot now has some context in the form of these titles and labels, and is far more understandable, as shown in the following screenshot:



Let's now add titles and axis labels to base plots and utilize the `main`, `sub`, `xlab`, and `ylab` options to change the titles and axis labels of base plots. Follow the steps given below:


1. Load the `iris` dataset using `data("iris")`.
2. Plot petal length and width from the `iris` dataset to see what the plot looks like, and take note of the default axis labels as follows:

```
plot(iris$Petal.Length, iris$Petal.Width)
```

3. Now, add a title, subtitle, and custom axis labels to the same plot using the following code:


```
plot(iris$Petal.Length, iris$Petal.Width,  
     main = "Petal Width vs. Length",  
     sub = "iris dataset",  
     xlab = "Petal Length",  
     ylab = "Petal Width")
```

Output: The following is the output we get when we execute the code line mentioned in [Step 2]:



The following is the output we get when we execute the code mentioned in [Step 3]:

Petal Width vs. Length



If we decide we'd like to plot in a different color, say red, it's as simple as passing `col = "red"` into `plot()`. R supports the names of many different colors along with hexadecimal color codes. The code to change the previous plot to red would be as follows:

```
plot(mtcars$wt, mtcars$mpg,
     main = "mpg vs. wt, mtcars data",
     xlab = "weight",
     ylab = "mpg",
     col = "red")
```

Changing the Color of Base Plots

Let's see how you can use the `col` option provided by the `plot()` function to change a plot into a few different colors. Follow the steps given below:

1. Use the `col` option to turn the plot from the last exercise blue as follows:

```
plot(iris$Petal.Length, iris$Petal.Width,
     main = "Petal Width vs. Length",
     sub = "iris dataset",
     xlab = "Petal Length",
     ylab = "Petal Width",
     col = "blue")
```

2. Use the `col` option to turn the plot from the last exercise yellow using the hexadecimal color code `"111111"`:

```
plot(iris$Petal.Length, iris$Petal.Width,
     main = "Petal Width vs. Length",
     sub = "iris dataset",
     xlab = "Petal Length",
     ylab = "Petal Width",
     col = "111111")
```

Output:

1. Check your `plot` window after executing the code in [Step 1] to be sure that the plot is now blue.
2. Check your `plot` window after executing the code in [Step 2] to be sure that the plot is now yellow.

It is important to know and understand `plot()`, as base plots are adequate and useful. However, the `ggplot2` package has really taken over the R graphics landscape, and as such we won't spend much more time on base plots. Let's do a quick activity just to be sure we have the hang of them.

Activity: Recreating Plots with Base Plot Methods

Scenario

You have been asked to create some base plots that provide information on the `mtcars` and `iris` datasets for a junior colleague.

Prerequisites


Make sure you have R and RStudio installed on your machine.

Aim

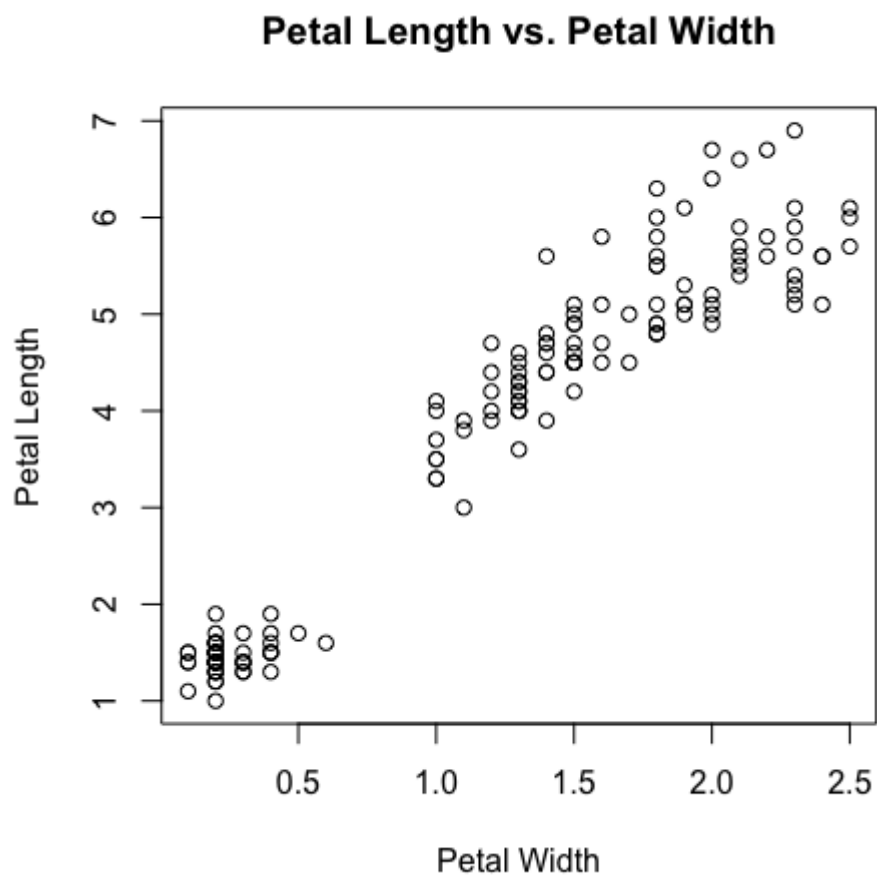
To use `plot()` by recreating different plots with different base plot methods.

Steps for completion

1. Load the datasets library using `library(datasets)`.
2. Load the `iris` and `mpg` datasets. You will need to make individual calls, using `data("mtcars")`, for example. You will then see the dataset in your environment as a promise. It will appear as a dataset in your list of datasets in the upper-right window when you first attempt to use it.
3. Recreate the following base plots using `iris` data:
 - a. A scatterplot to plot petal width without axis labels:





1. b. A scatterplot to plot petal length and width with axis labels:




1. c. Scatterplots in `[*[*1`

```
× *][[*2*] grids to plot petal  
length and width with axis labels:
```



4. Recreate the following histogram using `mtcars` data to plot the number of cylinders in the color blue:



ggplot2

ggplot2 is an incredibly popular graphics package in R. It can be installed on its own or comes as part of the `Tidyverse` set of packages.

When you're using ggplot2, both as you're learning how to use it and even when you're more seasoned, the official RStudio ggplot2 cheat sheet will be a resource you may want to keep close for your reference. It will not only remind you of the basics (and more advanced implementations) of how to use ggplot2, it gives suggestions for which plots to use when you have certain types of variables (for example, if you have one continuous variable, you can build a histogram using `geom_hist()`).

Note:

The ggplot2 cheat sheet can be found at the following URL: <https://www.rstudio.com/wp-content/uploads/2015/03/ggplot2-cheatsheet.pdf>.

RStudio has also made many different cheat sheets available for common R packages. They can be found on their official website at the following URL: <https://www.rstudio.com/wp-content/uploads/2015/03/ggplot2-cheatsheet.pdf>

First and foremost, you'll need to install ggplot2 using `install.packages("ggplot2")` or through point-and-click methods. Then, when you load ggplot2 in RStudio using `library(ggplot2)`, it immediately suggests the ggplot2 **Stack Overflow** tag as a good place to go for any help you might need ggplotting.



```
> library(ggplot2)
Stackoverflow is a great place to get help:
http://stackoverflow.com/tags/ggplot2.
```


Note:

The online documentation for ggplot at the [Tidyverse](http://ggplot2.tidyverse.org/index.html) website is thorough and can be used to supplement the built-in documentation. The URL is as follows: <http://ggplot2.tidyverse.org/index.html>. It contains many examples and thorough explanations of every element of ggplot2 and is maintained by the authors of the package.

ggplot2 Basics

To begin with, here is the exact same data, plotted both with `plot()` and `ggplot()`, respectively:





This is the built-in `cars` dataset, which contains only two variables, `speed` and `dist`. You can generate these plots yourself, as follows:

- Plot 1:

```
plot(cars)
```

- Plot 2:

```
library(ggplot2)
ggplot(cars, aes(speed, dist)) + geom_point()
```

Voilà! A plot and a ggplot. Which is more aesthetically pleasing to you? Which would you rather publish on a report or your blog? The answer is [*probably*] the ggplot, if you're like most data scientists out there.

Using ggplot2 requires you to begin to think of each element of a plot as a layer.

First, you have a white screen with only axes defined, two lines symbolizing the `[x]` and `[y]` axes. Using the `ggplot()` function, you layer on a dataset that contains what you'll plot and the aesthetics of the plot, defined in the `aes()` function, which corresponds to the things to plot and how to plot them. Then, you layer on a geom, using a `geom_*()` function, which tells ggplot2 what kind of plot you're trying to make. You can layer on additional aesthetics, such as plot titles, axis labels, colors, different point types, and more.

```
ggplot(data = <DATA>) + <GEOM_FUNCTION>(mapping = aes(<MAPPINGS>))
```

In the call, we see the three things required by all ggplots:

1. Dataset (`DATA`)
2. Geom (`GEOM_FUNCTION`)
3. Mappings (`MAPPINGS`)

Your dataset, entered in `DATA`, will be the dataset you're looking to plot variables from. Geoms take the form of `geom_*()`, where the `*` will be the name of the type of plot you're looking to create, for example `geom_point()` for a scatterplot, `geom_boxplot()` for a boxplot, and `geom_histogram()` for a histogram (perhaps you're detecting a theme here in how the `geom_*()` functions are named!)

Mappings are the variables you want to graph plus other aesthetics (`aes()` is short for aesthetics)

```
ggplot(data = <DATA>;, aes(<GLOBAL MAPPINGS>)) + <GEOM_FUNCTION>(  
  mapping = aes(<LOCAL MAPPINGS>))
```

Global mappings will apply globally to every layer of your plot. This is a good place to put the variables you'd like to begin plotting and any settings you'd like to apply to everything, for example declaring `alpha = 0.6` here would mean all of your points in a scatterplot are at 60% transparency.

Local mappings will either override or add to any global mappings and apply to that layer only. As you'll see later on, you can include a number of layers in your ggplot (either by plotting multiple variables or by adding layers that include titles or themes, for example), so any local mappings should be applied inside the `geom_*()` or appropriate function (for example, `ggtitle()`).

There are a few things you should know about creating ggplots that will help you along the way. Firstly, you can save a ggplot call and use it for multiple graphs, for example:

```
#save the ggplot data and mappings as 'mtcars_ggplot':  
mtcars_ggplot <- ggplot(mtcars, aes(wt, mpg))  
#create 2 additional plots:  
mtcars_ggplot + geom_point()  
mtcars_ggplot + geom_point(aes(col = factor(cyl)))
```

The first line of code saves a ggplot object called `mtcars_ggplot`, which says that you want to use the `mtcars` dataset and the weight (`wt`) and miles per gallon (`mpg`) variables for plotting. This object will be saved in your R environment as a list, and you can view it in the environment by hitting the magnifying glass icon next to its name:

Name	Type	Value
mtcars_ggplot	list [9] (S3: gg, ggplot)	List of length 9
data	list [32 x 11] (S3: data.frame)	A data.frame with 32 rows and 11 columns
layers	list [0]	List of length 0
scales	environment [1] (S3: ScalesLis	<environment: 0x108eda358>
mapping	list [2] (S3: uneval)	List of length 2
x	symbol	`wt`
y	symbol	`mpg`
theme	list [0]	List of length 0
coordinates	environment [3] (S3: CoordCa	<environment: 0x108ed8ae8>
facet	environment [3] (S3: FacetNul	<environment: 0x108ed76d8>
plot_env	environment [4]	<environment: R_GlobalEnv>
labels	list [2]	List of length 2
x	character [1]	'wt'
y	character [1]	'mpg'

You can see from inside the `mtcars_ggplot` object that it has saved the `mtcars` dataset in the data list. There are currently no layers, because we haven't told it what kind of plot we want yet (or any titles or axis labels). You can see in the mapping list that `[x]` is now the `wt` variable and `[y]` is the `mpg` variable, which we indicated by putting those variables in as mapping arguments inside `aes()`. The default labels, which are taken directly from the variable names as they are saved in the dataset, are also listed in the labels list, as `[x] = wt` and `[y] = mpg`, which are both character strings.


The second line of code plots a basic scatterplot of miles per gallon by weight of the car. It calls the `mtcars_ggplot` object for beginning guidance (`DATA` and `MAPPINGS`) and then uses `geom_point()` as its `GEOM`. The third line of code also creates a scatterplot, similarly calling `mtcars_ggplot` as its guide, but adds in an additional local mapping inside `geom_point()`, declaring that it wants the points colored (`col`) by the factor variable `cyl`, which indicates how many cylinders the car has. If you first load `mtcars` in your R environment, if you haven't already, using `data("mtcars")`, all of this code is executable in R. Feel free to try it to see the different plots.

Secondly, the plus signs you'll need to add layers to a `ggplot()` object must *[always]* come at the end of a line. The following code will run successfully to create the plot we saw at the beginning of this subtopic:

```
ggplot(cars, aes(speed, dist)) + geom_point()
```

The following code will not run, because the plus sign has been moved down to the second line, in front of `geom_point()`:

```
ggplot(cars, aes(speed, dist))
+ geom_point()
```



If you attempt to run code with the plus sign at the beginning of a line, preceding `geom_point()`, as in the previous example, a blank plot, as shown in the preceding screenshot, will generate in your **Plots** window in RStudio and you will get the following error in your console:

```
Error in +geom_point() : invalid argument to unary operator
```

Let's walk through some basic types of ggplots using the `mtcars` dataset, which we've used a few times so far and will continue to use throughout the course. Built-in datasets in R are convenient to use for trying to learn new things; plus they can be helpful to use for creating examples for others when we need help with something.


Histogram

When you have one continuous variable, it's a good idea to use a histogram to get an idea of its distribution. The height of the bar of the histogram corresponds to the number of observations that have that value. We can create a histogram of the `mpg` variable in `mtcars` using the following code:

```
ggplot(mtcars, aes(mpg)) + geom_histogram()
```


This code will throw a warning:

```
'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




The default number of bins is always 30, and you should always change it and find a better value for your data. The default means that it takes the range of the data (here, `mpg` is between 10.4 and 33.9) and divides it by 30 to create bins---which, in this case, is a bit large, and causes our binwidth to be equal to 0.783, which is tiny! Let's choose a few different bin widths to see what happens. Note that you need to specify `binwidth` inside of `geom_histogram()` as a local mapping.


`binwidth = 10` gives us almost no detail---we can see three groups of observations of the data:



`binwidth = 1` isn't bad, but the graph shows some gaps. Let's see if we can close them:



Using a binwidth of 3 shows decent amount of detail, as shown in the following graph:



Creating Histograms using ggplot2

In this section, we will create a histogram with `ggplot2` and experiment with different binwidths to find the best representation of the data. Follow the steps below:

1. Install the `ggplot2` library and then load it:

```
install.packages("ggplot2")  
library(ggplot2)
```

2. Load the `msleep` dataset, a built-in dataset that comes installed with `ggplot2`, using `data("msleep")`.
3. Create a histogram of the `sleep_total` variable from `msleep`. Do you get the binwidth error?

```
ggplot(msleep, aes(sleep_total)) + geom_histogram()
```


4. Try the same histogram, but with `binwidth = 10`. Does the histogram improve?

```
ggplot(msleep, aes(sleep_total)) + geom_histogram(binwidth = 10)
```


5. Try the histogram one more time, now with `binwidth = 1`:

```
ggplot(msleep, aes(sleep_total)) + geom_histogram(binwidth = 1)
```


Output: We get a binwidth error along with the following graph when we try to create a histogram of the `sleep_total` variable from the `msleep` dataset using the code mentioned in [Step 3]:



The following is the histogram with `binwidth = 10` :




The following is the histogram with `binwidth = 1`:



Bar Chart

For one categorical or factor variable, you can create a bar chart. We can create a bar chart of the `cyl` variable of `mtcars` using the following code:


```
#using geom_bar()
ggplot(mtcars, aes(cyl)) + geom_bar()
```



One fun fact is that we can actually create bar charts with the `geom_histogram()` call as well, by including `stat = "count"`, as follows:

```
#using geom_histogram() and stat
ggplot(mtcars, aes(cyl)) + geom_histogram(stat = "count")
```

You can ignore the warning it will throw; this creates the exact same bar chart:



Creating a Bar Chart with ggplot2 using Two Different Methods

Let's create a bar chart with ggplot2 using both `geom_bar()` and `geom_hist()`.


1. Create a bar chart of the `vore` variable from `msleep` using `geom_bar()`, as follows:

```
ggplot(msleep, aes(vore)) + geom_bar()
```


2. Create the same bar chart of the `vore` variable from `msleep` using `geom_histogram(stat = "count")`, as follows:

```
ggplot(msleep, aes(vore)) + geom_histogram(stat = "count")
```

Output: The following is the output we get after executing the code mentioned in [Step 1]:



The following is the output we get after executing the code mentioned in *[Step 2]*:




Scatterplot

Two continuous variables are a good candidate for a scatterplot, which is created using `geom_point()` in ggplot2. You can also create scatterplots using `geom_jitter()`. Using `jitter` instead of `point` adds a small amount of noise (tiny on the order of decimals) to each observation, spreading them out from one another slightly so they're easier to see.

We can create a scatterplot with the `wt` and `mpg` variables from `mtcars` using the following code:

```
ggplot(mtcars, aes(wt, mpg)) + geom_point()
```





The plot shows a clear relationship between the weight and miles per gallon of the cars, namely as `wt` increases, the `mpg` decreases (we saw this same relationship when we created this plot as a base plot earlier in this lab).

Though we won't really be able to see much of an effect with this dataset, you can create a scatterplot with a bit of jitter introduced using the following code:

```
ggplot(mtcars, aes(wt, mpg)) + geom_jitter(width = 0.1)
```

You can control exactly how much jitter by inputting `width =` some number into `geom_jitter()` :



Creating a Scatterplot of Two Continuous Variables

Let's now create scatterplots using `geom_point()` and `geom_jitter()`.


1. Create a scatterplot of the `bodywt` and `sleep_total` variables from `msleep`:

```
ggplot(msleep, aes(bodywt, sleep_total)) + geom_point()
```


2. This scatterplot is a great candidate for using `geom_jitter()`, as many of the `sleep_total` observations cluster around the zero bodyweight. We'll use a fairly large width for jitter to really separate these points, because the scale of `bodywt` is in the thousands:

```
ggplot(msleep, aes(bodywt, sleep_total)) + geom_jitter(width = 50)
```

Output: The following is the output we get when we execute the code mentioned in [Step 1]:



The following is the output we get when we execute the code mentioned in [Step 2]:



Boxplot

Boxplots are most appropriate when you want to check the distribution of a continuous [y] variable with some categorical (factor) [x] variable. We cannot create a boxplot of the `mpg` variable with the `cyl` variable in `mtcars` using the following code:

```
ggplot(mtcars, aes(cyl, mpg)) + geom_boxplot()
```


We will get a warning as follows:

```
Warning message: Continuous x aesthetic -- did you forget aes(group=...)?
```

The `cyl` variable is not explicitly declared as a factor variable in the `mtcars` dataset, so `ggplot` is confused about what the [x] variable is supposed to be. This is similar to when we created base plots with factor variables, though as we saw, `plot()` will still plot a variable not declared as a factor, but it will create a scatterplot instead of the desired histogram. The following code, which transforms `cyl` into a factor variable using `as.factor()`, will fix it and plot the boxplot correctly:

```
ggplot(mtcars, aes(as.factor(cyl), mpg)) + geom_boxplot()
```

Thus, we get the following graph as an output:



Of course, now the axis label reads `as.factor(cyl)`, because the default axis label is whatever is input as `[x]`. We'll learn how to fix that in the next subtopic!

Creating Boxplots using ggplot2

Let us create a boxplot using `geom_boxplot()`.


1. Create a boxplot of `sleep_total` with `vore`, both variables from the `msleep` dataset using the following code:

Note:

Notice that `omni` seems to have four outliers, represented by the black dots outside of the boxes, which represent the **Interquartile Range (IQR)** of the `sleep_total` of each variable.

```
ggplot(msleep, aes(vore, sleep_total)) + geom_boxplot()
```

Output: We get the following boxplot as the output after executing the preceding code:



While these four types of plots are far from everything available in `ggplot2`, everything we've gone over so far in this subtopic should be enough to get started creating basic `ggplots`. To start, we should get comfortable with building the basics, and then we'll extend them using other calls to `aes()`, plus titles and custom axis labels.

Activity: Recreating Plots Using `ggplot2`

Scenario

You have been asked to create some `ggplots` that provide information on the `mtcars` and `iris` datasets for a presentation in your office.

Prerequisites

You should have RStudio and R installed on your machine. The `ggplot2` package should also be installed.


Aim

To construct basic `ggplots` by recreating some of those shown in the preceding exercises.


Steps for Completion

1. Load `ggplot2` using `library(ggplot2)` .
2. Try to recreate all of the following `ggplots` using the `iris` dataset:

a. A histogram to plot petal width:




1. b. A scatterplot to plot petal length and width:




1. c. Boxplot to plot petal width and the

```
`Species` factor variable:
```

4. Try to recreate the following bar chart ggplot using the `gear` variable of the `mtcars` dataset:



Digging into aes()


While we have created some basic ggplots, we haven't really dug much into the aesthetics of plots. There are definitely both some global and plot-specific aesthetics that are very important to know when you're building plots.

One key distinction to master is that when you call something inside of `aes()`, the aesthetic is mapped to the value of the variable in the data. Outside of an `aes()` call, the aesthetic is set to a specific value. This is perhaps best understood with an example.

The following code is a bar chart of how many cars have each number of cylinders, where `fill` is the number of gears the car has, all from `mtcars`:

```
ggplot(mtcars, aes(cyl, fill = as.factor(gear))) + geom_bar()
```

A legend appears to let us know which color corresponds to which number of gears, as shown in the following graph:



The `fill` is inside of `aes()` and the variable is entered as a factor, both of which are required for this to work.

If, instead, you were looking to make all of the bars light blue, seeing the preceding code, you might be tempted to run the following code:

```
ggplot(mtcars, aes(cyl, fill = "lightblue")) + geom_bar()
```


Or even the following code:

```
ggplot(mtcars, aes(cyl, fill = lightblue)) + geom_bar()
```

However, this code is looking for a thing called `lightblue` in the dataset, because you entered it inside of `aes()`. To actually fill the bars light blue, you should use:

```
ggplot(mtcars, aes(cyl)) + geom_bar(fill = "lightblue")
```

This produces the following graph of the bar chart of the count of cars of each cylinder type, but the bars have been colored light blue:



There are some very helpful global and local options that you'll probably need when you're using ggplot2 to create different plots. Let's go through a few of them.

Bar Chart

To make these charts better, we're going to convert the `cyl` and `gear` variables in `mtcars` to factor variables using the following code:

```
mtcars$cylfactor <- as.factor(mtcars$cyl)
mtcars$gearfactor <- as.factor(mtcars$gear)
```

Use of the factor variables will help the data display properly.

We previously saw how to both automatically change the color of a bar chart (when we made them light blue) and also how to fill a bar chart with another variable. We did this using the following code:


```
ggplot(mtcars, aes(cyl, fill = gearfactor)) + geom_bar()
```

The `fill` indicates the count of each car with a particular type of cylinder and gear. There are a few other ways to display the fill that we can use.

If we want the bars to be next to each other instead, we can add `position = "dodge"` inside `geom_bar()`, with the following code:

```
ggplot(mtcars, aes(cyl, fill = gearfactor)) + geom_bar(position = "dodge")
```

The output will be as shown in the following screenshot:




Now, the bars are all next to each other, and it's actually somewhat easier to see that there are no eight-cylinder cars with four gears. ggplot2 also automatically adds a legend to ggplots when you're plotting with colors or shapes.

If we want the bars to reflect percentages instead of representing the count of cars with a certain number of gears and cylinders, we can add `position = "fill"` inside `geom_bar()` :

```
ggplot(mtcars, aes(cyl, fill = gearfactor)) + geom_bar(position = "fill")
```

The output we get is as follows:



While the [x]-axis still says `count` (we'll be rid of this soon!), it has rescaled from 0 - 1.00, because it represents the percentages instead of counts.

Using Different Bar Chart Aesthetic Options

Let's now create bar charts using the `dodge` and `fill` bar chart position aesthetic options.


1. If not loaded from the last topic, load the `msleep` dataset using `data("msleep")`.
2. Create a bar chart using the `dodge` position aesthetic of `vore`, filled with the `conservation` variable. (These variables are already declared as factor variables when you load `msleep`.) The code for this is as follows:

```
ggplot(msleep, aes(vore, fill = conservation)) + geom_bar(position = "dodge")
```


3. Create a bar chart with the same variables, this time using the `fill` position aesthetic:

```
ggplot(msleep, aes(vore, fill = conservation)) + geom_bar(position = "fill")
```

Output: The following is the output we get when we execute the code mentioned in [Step 1]:



The following is the output we get when we execute the code mentioned in [Step 2]:



Facet Wrapping and Gridding


Facet wrapping and gridding can be applied to any ggplot, not just bar charts. Facet wrapping will split the base ggplot (which, here, is the count of cars with each number of cylinders) by a second variable, which, here, will be the number of gears, generating three plots. The code for this is as follows:

```
ggplot(mtcars, aes(cylfactor)) + geom_bar() + facet_wrap(~gear)
```

We can see that each of the three numbers of gears (3 , 4 , 5) have a bar chart for the count of the number of cars with each of the three types of cylinders (here, `cylfactor`, with values 4 , 6 , 8). Facet wrapping can be applied to any of the ggplots, though it may sometimes look strange, which can be mitigated with facet gridding.

Facet gridding is closely related to facet wrapping but allows for gridding by (`row ~ column`). The following code will generate the same as the preceding facet wrapping code, as gear is in the column place:


```
ggplot(mtcars, aes(cylfactor)) + geom_bar() + facet_wrap(~gear)
```

However, if you move gear to the row place, it will grid the plots by row instead of column, as shown in the following code:

```
ggplot(mtcars, aes(cylfactor)) + geom_bar() + facet_grid(gear~)
```

Thus, the output we get will be as follows:



Note:

You must remember to put the period in the columns place, which stands for *[all columns]*, or the code will not run.

Utilizing Facet Wrapping and Gridding to Visualize Data Effectively

Let's create bar charts using the `facet_wrap()` and `facet_grid()` functions. Follow the steps below:


1. Create a bar chart of `conservation`, facet wrapped by `vore`, both variables from the `msleep` dataset, as shown in the following code:

```
ggplot(msleep, aes(conservation)) + geom_bar() + facet_wrap(~vore)
```


2. Create the same bar chart, but use `facet_grid()` to grid the charts by row instead of column, as shown in the following code:

```
ggplot(msleep, aes(conservation)) + geom_bar() + facet_grid(vore~)
```

Output: The following is the output we get when we execute the code in *[Step 1]*:



The following is the output we get when we execute the code in [Step 2]:




Boxplot + coord_flip()

One handy feature to know about is an additional aesthetic layer, `coord_flip()`. Given that R functions are named in an informative way, it probably does more or less exactly what you'd think.

Let's return to our boxplot example from `mtcars`, which shows the distribution of `mpg` by the number of cylinders. We modify the code and add `coord_flip()` as follows:

```
ggplot(mtcars, aes(cylfactor, mpg)) + geom_boxplot() + coord_flip()
```

The output we get will be as shown in the following screenshot:



We see that `cylfactor` is now on the [y]-axis and `mpg` is on the [x]-axis, and the boxplots can flip.

`coord_flip()` can be implemented on other ggplots as well, but boxplots are a good way to easily see its effect.

Scatterplot

Scatterplot requires a bit more care than the other graphs we've covered to be truly meaningful and visually appealing. We'll return to our `mtcars` example of plotting `mpg` versus `wt`. If you recall the code used to build boxplots, you might be tempted to color the scatterplots using `fill = cylfactor` and code that looks like this:

```
ggplot(mtcars, aes(wt, mpg, fill = cylfactor)) + geom_point()
```

When we run this, we see that a legend has appeared that has the appropriate values of `cylfactor`, but the dots are all still black, so we've gained no new information. How do you think we fix this?


If you thought to yourself [*We need to use `col = cylfactor` in that `aes()` call*], then you're absolutely right. The code should be:

```
ggplot(mtcars, aes(wt, mpg, col = cylfactor)) + geom_point()
```

The code will also run if you spell out `col` as `color`. We can change the shape of the points inside a scatterplot with the shape aesthetic. `shape = 17` makes all the points into little triangles. The code is as follows:

```
ggplot(mtcars, aes(wt, mpg)) + geom_point(shape = 17)
```

The output we get is as follows:



However, these are very tiny. Let's make them bigger with the size options, which we'll also specify inside of `geom_point()` itself. Here is the code for it:

```
ggplot(mtcars, aes(wt, mpg)) + geom_point(shape = 17, size = 3)
```

Much better! If we had specified color as well, these settings would also apply.

With the bigger size, there's a small amount of overlap with some of the triangles. We can control the transparency using `alpha = some number between 0 (transparent) and 1 (opaque)`. Let's start with `alpha = 0.6` and then adjust as needed. Here is the code for that:

```
ggplot(mtcars, aes(wt, mpg)) + geom_point(shape = 17, size = 3, alpha = 0.6)
```

Utilizing Different Aesthetics for Scatterplots, Including Shapes, Colors, and Transparencies

Let us create scatterplots and change the shape and size of points, the colors, and the transparency of points. Follow the steps given below:

1. To make these scatterplots more visually appealing, load `dplyr` and remove the two observations with a `bodywt` greater than 2000, creating the `msleep2` dataset, as follows:

```
library(dplyr) msleep2 <- msleep %>% filter(bodywt < 2000)
```


2. Now, create a scatterplot of `bodywt` versus `brainwt`, using triangles for the points. You will see an error in your console window saying that it removed rows with missing values. Don't worry about this for now; missing data isn't the focus of this exercise.

```
ggplot(msleep2, aes(brainwt, bodywt)) + geom_point(shape = 17)
```


3. Create the same scatterplot but make the triangles much bigger.

```
ggplot(msleep2, aes(brainwt, bodywt)) + geom_point(shape = 17, size = 6)
```

Output: The following is the output we get after executing the code mentioned in [Step 2]:



The following is the output we get after executing the code mentioned in [Step 3]:



We've covered some global and local aesthetics that are very useful when building ggplots. Let's do a few examples to help master them.

Activity: Utilizing ggplot2 Aesthetics

Scenario

You have been asked to create some ggplots that provide information on the `mtcars` and `iris` datasets for a presentation for your colleague, as shown in the following screenshot:

0	1	2	3	4	
□	○	△	+	×	
5	6	7	8	9	
◇	▽	⊠	✱	⬠	
10	11	12	13	14	
⊕	⊗	⊞	⊗	⊞	
15	16	17	18	19	
■	●	▲	◆	●	
20	21	22	23	24	25
●	●	■	◆	▲	▼

Aim


To get the students comfortable with using more aesthetic options in their ggplots by having them recreate a few, as shown.


Prerequisites

Make sure you have R and RStudio installed on your machine. The ggplot2 package should also be installed.


Steps for Completion


1. Load `ggplot2` using `library(ggplot2)` .
2. Try to recreate the ggplots shown as follows.
3. The plots use the following datasets:
 - Plots 1 and 2: `mpg`





1.
 - Plots 3 and 4: diamonds





Extending the Plots with Titles, Axis Labels, and Themes

Thus far, we've learned much of how ggplot2 works by creating a few often-used plots. We have not addressed adding titles, customizing axis labels, or themes. Let's look at the basics of these three things.

Titles and axis labels can be added in two different ways in ggplot2.

If you're interested in changing everything in one go, you can use the `labs()` function. `labs()` takes as an argument anything you'd like to change the label of, such as title, subtitle, [x], [y], caption, or even the label of the legend. It is often used as follows:

```
mtcars_ggplot + geom_point() +
  labs(title = "mpg vs. wt",
        subtitle = "mtcars dataset",
        x = "weight",
        caption = "decreasing linear trend")
```

This adds a title and a subtitle, changes the [x]-axis label, and adds a caption.


These changes can all be made individually with the corresponding individual functions, such as `ggtitle()`, `xlab()`, and `ylab()`. They can all be added individually as layers to your plot to adjust the corresponding aspects.

Themes make more sense when demonstrated. We'll only cover a few of the built-in ggplot2 themes, but know that building your own custom themes for ggplot2 is entirely possible.


Let's return to an example we've used a few times in this lab and see how a few different themes change the look. The `mtcars` dataset with `mpg` versus `wt`, colored by cylinder, is a great one. The default theme is `theme_gray()`, so we'll skip that one.

The other themes available are as follows:


- `theme_bw()` removes the gray background and makes it black and white. Very straightforward.
- `theme_classic()` removes everything behind the points: there are no secondary axis marks or fill, only blank white space, as shown in the following screenshot:



- `theme_dark()` makes everything behind the points a much darker gray. You can see that it also changes the legend to match, as shown in the following screenshot:



- `theme_minimal()` leaves the axis marks very light gray, but includes no fill. The output is as follows:



These are just a few theme examples. There are more listed in the documentation that you could experiment with, and you can create your own custom themes once you have more experience with ggplot.

Note:

When you're looking to save your ggplots, you have two options. You can use the **export** button above your plot on the **Plots** tab in the viewer, or you can use the `ggsave()` function.

`ggsave("my_mtcars_plot.png")` will save your plot as a PNG with the filename `my_mtcars_plot` in your working directory. If you want to specify another directory, you can do that as well, using `ggsave("images/my_mtcars_plot.png")`, which saves the plot instead in a folder called `images`.

Let us now add titles and axis labels to ggplots by extending the aesthetic options. Follow the steps given below:


1. Load `ggplot2` using `library(ggplot2)`.
2. Try to recreate the ggplots shown as follows. :
 - Execute the following code:

```
library(ggplot2)
ggplot(cars, aes(speed, dist)) + geom_point() +
labs(title = "dist vs. speed",
```



```
subtitle = "cars dataset",  
y = "distance")
```

The output of the preceding code will be as follows:




1.

- Execute the following code:

```
ggplot(cars, aes(speed, dist)) + geom_point() +  
ggtitle("dist vs. speed",  
        subtitle = "cars dataset") +  
ylab("distance")
```

The output of the preceding code will be as follows:




1.

- Execute the following code:

```
ggplot(diamonds, aes(carat, price, col = cut)) + geom_point(alpha = 0.4) +  
theme_minimal()
```

The output of the preceding code will be as follows:



Summary

Graphing in R will be crucial in your data science work, and we have covered most of the basics here. However, graphing is one of those things where, most of the time, there are always going to be different types of graphs you haven't heard of yet and options you haven't yet selected, so it's important to know where to look for assistance and how to keep learning.

Let's press forward on to the next topic, where we'll begin to look at some data more closely, doing some cleaning and data management necessary to get us one step closer to modeling and analysis.